

Wasil Engel
12231558
PSet 3

I collaborated with Gina Son.

To see more screenshots of the plots I created for the coding part of the assignment, please see my Jupyter notebook on Github here: <https://github.com/Wasil-UChi/Machine-Learning>

Ch 5 RESAMPLING METHODS

6

p.199

```
path = '/Users/wasilengel/Desktop/School/Harris/Machine Learning/Pset 3/Data-Default.csv'
```

```
default = pd.read_csv(path)
```

```
default.head() # 10000 rows × 5 columns
```

```
# a
```

```
encoding_dict = {"Yes": 1, "No": 0}
default["default"] = default["default"].map(encoding_dict)
default.head()
```

```
X = default[["balance", "income"]]
X = sm.add_constant(X)
y = default["default"]
```

```
results = sm.Logit(y, X).fit()
print(results.summary())
```

```
# The SE for balance are 0.000 and for income 4.99e-06.
```

```
# b
```

```
def get_indices(data, num_samples):
    np.random.seed(1)
    positive_data = data[data["default"] == 1]
    negative_data = data[data["default"] == 0]
```

```

    positive_indices = np.random.choice(positive_data.index, int(num_samples / 4), replace =
True)
    negative_indices = np.random.choice(negative_data.index, int(3*num_samples / 4), replace
= True)
    total = np.concatenate([positive_indices, negative_indices])
    np.random.shuffle(total)
    return total

```

```

def boot_fn(data, index):
    np.random.seed(1)
    X = data[["balance", "income"]].loc[index]
    y = data["default"].loc[index]
    lr = LogisticRegression()
    lr.fit(X, y)
    intercept = lr.intercept_
    coef_balance = lr.coef_[0][0]
    coef_income = lr.coef_[0][1]
    return [intercept, coef_balance, coef_income]

```

```

intercept, coef_balance, coef_income = boot_fn(default, get_indices(default, 100))
print(f"Intercept is {intercept}, the coeff of balance is {coef_balance}, the coeff for income is
{coef_income}")

```

c (purpose: compare SE -- however, unable because Python)

```

def boot(data, func, R):
    total_coeff_balance = []
    total_coeff_income = []
    for i in range(R):
        bootstrap = resample(data, replace=True, n_samples=(0.3*default.size), random_state=i,
stratify = data['default'])
        params = func(data, bootstrap.index)
        total_coeff_balance.append(params[0])
        total_coeff_income.append(params[1])
    return (np.mean(total_coeff_balance), np.mean(total_coeff_income))

```

```

total_coeff_balance, total_coeff_income = boot(default, boot_fn, 1000)

```

```

print(f"The total coeff of balance is {total_coeff_balance} and the total coeff of income is
{total_coeff_income}")

```

d

Since this is a simplified version of the R boot function, I am unable

to get the standard errors, however, I would expect them to be very

```
# similar -- unlike the coefficient estimates/ means, which are at  
# -6.995 for balance and at 0.004 for income and thus different from  
# the coefficients in part a) which vary unless I set seed in boot_fn().
```

8

p.200f.


a

```
np.random.seed(1)  
y = np.random.normal(size = 100)  
X = np.random.normal(size = 100)  
y = X - 2 * (X ** 2) + np.random.normal(size = 100)
```

```
# n is 100 -> number of observations/ units  
# p is 2 -> number of independent variables/ features (X, X^2)  
# equation is  $y = X - 2X^2 + e$  // where e = error term
```

b

```
sns.scatterplot(X, y)  
plt.xlabel("X")  
plt.ylabel("y")  
plt.title("scatter plot")
```



```
# Comment on what you find: simulated resembles a quadratic function that  
# ranges from -15 to 2.5 along the y-axis and -2 to 2 along the x-axis.
```

c

```
np.random.seed(1)  
  
for i in range(1, 5):  
    poly = PolynomialFeatures(i, include_bias = False)
```

```

predictors = poly.fit_transform(X.reshape(-1, 1))
lr = LinearRegression()
error = -1 * cross_val_score(lr, predictors, y, cv = len(X), scoring =
"neg_mean_squared_error").mean()
print(f"For model {i}, error is {error}")

```

d

```

np.random.seed(50374)
for i in range(1, 5):
    poly = PolynomialFeatures(i, include_bias = False)
    predictors = poly.fit_transform(X.reshape(-1, 1))
    lr = LinearRegression()
    error = -1 * cross_val_score(lr, predictors, y, cv = len(X), scoring =
"neg_mean_squared_error").mean()
    print(f"For model {i}, error is {error}")

```

Our results are identical because in LOOCV, there is no random sampling,
instead, it is being trained on the same n-1 observations/ folds (and
then just tested on the observations that is being left-out).

e

Yes, I expected the second model to have the smallest LOOCV test error
because of its underlying quadratic form, which we have seen in b (see
scatterplot). This is what best fits the second model, which is
the one with a quadratic form too: $Y = \beta_0 + \beta_1X + \beta_2X^2 + e$.

f

```

for i in range(1, 5):
    poly = PolynomialFeatures(i)
    predictors = poly.fit_transform(X.reshape(-1, 1))
    results = sm.OLS(y, predictors).fit()
    print(results.summary())

```

When I run the four model in a OLS regression, the coefficient estimates
for x1 is statistically significant for the first/ linear model(0.003).
Beyond that, the x1 and x2 estimates are even more highly significant
for the second (quadratic) model, third, and fourth models (at a level
of 0.000). Note how neither x3 in model 3 nor x3 and x4 are significant
at conventional significance levels. This suggests that our results
agree with our previous findings from CV where especially the quadratic
function yielded the smallest LOOCV test error (see answer to part e).

Ch 6 LINEAR MODEL SELECTION & REGULARIZATION

11

p.264

```
# a

## Explore

boston = load_boston()

data = pd.DataFrame(boston.data, columns=boston.feature_names)
data.head()

predictors = data.drop("CRIM", axis = 1)
X = data.drop("CRIM", axis = 1)
y = data["CRIM"]
predictors.head()

for feature in predictors.columns:
    sns.scatterplot(predictors[feature], data["CRIM"])
    plt.title(feature)
    plt.show()

## Best subset selection

# Quite frankly, this does not add up for me. Why are we hand-selecting
# features for best subset selection when our outcome should be a subset
# of variables that best suit our model? Hmm. Instead, we calculate the
# error rate here for a model that we choose based on visual correlation?
# But how is that best subset selection as demonstrated in Friday's lab?

hand_selected_features = ["NOX", "DIS", "RAD", "LSTAT"]
# Based on explored/ visualized data correlation from scatterplots

results_dict = {}

lin_reg = LinearRegression()
error = cross_val_score(lin_reg, predictors[hand_selected_features], y, cv=5, scoring =
"neg_mean_squared_error")
print("Error for best subset selection is", -np.mean(error))
results_dict["Best_subset"] = -np.mean(error)

## Forward stepwise selection
```

Option 1:

```
P = len(X.columns)
used_pred = []
M = []
M_scores = []

for K in range(P):
    best_score = -1000
    best_pred = None

    # Inner loop
    for var in X.columns:

        # Skips if predictor already used
        if var not in used_pred:
            predictors = used_pred[:]
            predictors.append(var)

            score = np.mean(cross_val_score(lin_reg, X[predictors], y, cv = 5, scoring =
'neg_mean_squared_error'))
            if score > best_score:
                best_score = score
                best_pred = var

        # Updates the list of used predictors and list of Mk models
        used_pred.append(best_pred)
        M.append(used_pred[:])
        M_scores.append(best_score)

best_M = M_scores.index(max(M_scores))
print('Predictors that make the best model are: ', M[best_M])
```

Option 2:

```
sfs1 = SequentialFeatureSelector(lin_reg,
                                k_features="best",
                                forward=True,
                                scoring='neg_mean_squared_error',
                                cv=5)

sfs1.fit(X, y)
sfs1.k_feature_names_
```

Backwards stepwise selection

```
sfs1 = SequentialFeatureSelector(lin_reg,
                                k_features="best",
                                forward=False,
                                scoring='neg_mean_squared_error',
                                cv=5)
```

```
sfs1.fit(X, y)
sfs1.k_feature_names_
```

```
# In the forwards subset selection, we start out with no Xs and add
# those X that most fit and repeat the process until all Xs are in the
# model. In the backwards subset selection, we start with all Xs and
# then remove X's that least fit and repeat the process until there are
# no Xs left in the model. Doing that,
# 1. forwards selection yields: ['RAD', 'LSTAT', 'ZN'] -- same by the
# way for both of the options I explored, while
# 2. backward selection yields: ('ZN', 'NOX', 'DIS', 'RAD', 'LSTAT').
# Moving forward, I'll use only those three variables that appear to
# contribute to best model fit resulting from both methods, namely:
# ['RAD', 'LSTAT', 'ZN']
```

```
# Going back to best subset selection, as I said before, what we
# learned in lab on how to do it just does not add up for me (cf.
# above): why are we hand-selecting features when our outcome itself
# should be a subset of variables that render best fit?
```

```
# b
```

```
# Evaluate model performances for each of the Xs I got in a)
# using K-Fold-CV -> so that is for: ['RAD', 'LSTAT', 'ZN']
```

```
error_list = []
for power in range(1, 11):
    X = data['RAD']
    y = data["CRIM"]
    poly = PolynomialFeatures(power, include_bias = False)
    X = poly.fit_transform(X.to_frame())
    lr = LinearRegression()
    error_list.append(-1*cross_val_score(lr, X, y, cv=10,
scoring="neg_mean_squared_error").mean())
print("K Fold CV")
print('RAD')
mini = min(error_list)
print(f"min MSE is: {mini}")
```

```
pd.DataFrame({"Degree": np.arange(1,11), "CV Mean Squared Error": error_list})
```

```
error_list = []
for power in range(1, 11):
    X = data['LSTAT']
    y = data["CRIM"]
    poly = PolynomialFeatures(power, include_bias = False)
    X = poly.fit_transform(X.to_frame())
    lr = LinearRegression()
    error_list.append(-1*cross_val_score(lr, X, y, cv=10,
scoring="neg_mean_squared_error").mean())
print("K Fold CV")
print('LSTAT')
mini = min(error_list)
print(f"min MSE is: {mini}")
pd.DataFrame({"Degree": np.arange(1,11), "CV Mean Squared Error": error_list})
```

```
error_list = []
for power in range(1, 11):
    X = data['ZN']
    y = data["CRIM"]
    poly = PolynomialFeatures(power, include_bias = False)
    X = poly.fit_transform(X.to_frame())
    lr = LinearRegression()
    error_list.append(-1*cross_val_score(lr, X, y, cv=10,
scoring="neg_mean_squared_error").mean())
print("K Fold CV")
print('ZN')
mini = min(error_list)
print(f"min MSE is: {mini}")
pd.DataFrame({"Degree": np.arange(1,11), "CV Mean Squared Error": error_list})
```

Judging by the lowest CV MSE which I printed out on top of each table
for the variables I previously identified in part a), I would propose
a model with 2 degrees because that yields the lowest *TOTAL* CV MSE.

The reason is that even though ZN and RAD both suggest that 7 is the
model resulting in the lowest CV MSE, we need to take the sum of total
test errors into account should we do that. Note how for ZN and RAD
the MSE rates are pretty much constant for all polynomial model degrees.
They differ by less than one percentage point for the degree values we
care about, therefore it doesn't matter to much for them what model
degree we pick. This is very different for LSTAT. Here, the CV MSE
differs notably across the degrees which is why I prioritize choosing
its lowest CV MSE because that returns the lowest CV MSE for the model


```

# overall. Since it's lowest for 2 degrees for LSTAT, I shall use a
# quadratic model function.

# c

# The forwards and backward subset selection methods both return a subset
# of the variables which is what the name suggests too: # ['RAD', 'LSTAT',
# 'ZN']. As I already elaborated on in greater detail in part a):

# "In the forwards subset selection, we start out with no Xs and add
# those X that most fit and repeat the process until all Xs are in the
# model. In the backwards subset selection, we start with all Xs and
# then remove X's that least fit and repeat the process until there are
# no Xs left in the model. Doing that,
# 1. forwards selection yields: ['RAD', 'LSTAT', 'ZN'] -- same by the
# way for both of the options I explored, while
# 2. backward selection yields: ('ZN', 'NOX', 'DIS', 'RAD', 'LSTAT').
# Moving forward, I'll use only those three variables that appear to
# contribute to best model fit resulting from both methods, namely:
# ['RAD', 'LSTAT', 'ZN']"

```